

Gaussian Processes - Computer Lab

Deadline: See LISAM
Teacher: Mattias Villani
Grades: Pass/Fail
Submission: Via LISAM

You should use R to solve the lab since the computer exam will be in R.
Output: an individual report and a group report to be presented on the seminar.
Attach your code in LISAM as separate files.

1. **Implementing Gaussian process regression from scratch.** This first exercise will have you writing your own code for the Gaussian process regression model:

$$y = f(x) + \varepsilon \quad \varepsilon \sim N(0, \sigma_n^2) \\ f \sim GP[0, k(x, x')].$$

When it comes to the posterior distribution for f , I **strongly** suggest that you implement Algorithm 2.1 on page 19 of Rasmussen and Williams (RW) book. That algorithm uses the Cholesky decomposition (`chol()` in R) to attain numerical stability. Here is what you need to do:

- (a) Write your own code for simulating from the posterior distribution of $f(x)$ using the squared exponential kernel. The function (name it `posteriorGP`) should return vectors with the posterior mean and variance of f , both evaluated at a set of x -values (x^*). You can assume that the prior mean of f is zero for all x . The function should have the following inputs:
- i. `x` (vector of training inputs)
 - ii. `y` (vector of training targets/outputs)
 - iii. `xStar` (vector of inputs where the posterior distribution is evaluated, i.e. x^* . As in my slides).
 - iv. `hyperParam` (vector with two elements σ_f and ℓ)
 - v. `sigmaNoise` (σ_n).
- [Hint: I would write a separate function for the Kernel (see my `GaussianProcess.R` function on the course web page) which is then called from the `posteriorGP` function.]
- (b) Now let the prior hyperparameters be $\sigma_f = 1, \ell = 0.3$. Update this prior with a single observation: $(x, y) = (0.4, 0.719)$. Assume that the noise standard deviation is known to be $\sigma_n = 0.1$. Plot the posterior mean of f over the interval $x \in [-1, 1]$. Plot also 95% probability (pointwise) bands for f .

x	-1.0	-0.6	-0.2	0.4	0.8
y	0.768	-0.044	-0.940	0.719	-0.664

Table 1: Simple data set for GP regression.

- (c) Update your posterior from 1b) with another observation: $(x, y) = (-0.6, -0.044)$. Plot the posterior mean of f over the interval $x \in [-1, 1]$. Plot also 95% probability bands for f . [Hint: updating the posterior after one observation with a new observation gives the same result as updating the prior directly with the two observations. Bayes is beautiful!]
- (d) Compute the posterior distribution of f using all 5 data points in Table 1 below (note that the two previous observations are included in the table). Plot the posterior mean of f over the interval $x \in [-1, 1]$. Plot also 95% probability intervals for f .
- (e) Repeat 1d), this time with the hyperparameters $\sigma_f = 1, \ell = 1$. Compare the results.

2. **Gaussian process regression on real data using the kernlab package.** This exercise lets you explore the kernlab package on a data set of daily mean temperature in Stockholm (Tullinge) during the period January 1, 2010 - December 31, 2015. I have removed the leap year day February 29, 2012 to make your life simpler. You can read the dataset with the command:

```
read.csv('https://github.com/STIMALiU/AdvMLCourse/raw/master/GaussianProcess/Code/TempTullinge.csv', header=TRUE, sep=',')
```

Create the variable `time` which records the day number since the start of the dataset (i.e. `time = 1, 2, ..., 365 \cdot 6 = 2190`). Also, create the variable `day` that records the day number since the start of each year (i.e. `day = 1, 2, ..., 365, 1, 2, ..., 365`). Estimating a GP on 2190 observations can take some time on slower computers, so let's thin out the data by only using every fifth observation. This means that your time variable is now `time = 1, 6, 11, ..., 2186` and `day = 1, 6, 11, ..., 361, 1, 6, ..., 361`.

- (a) Familiarize yourself with the following functions in `kernlab`, in particular the `gausspr` and `kernelMatrix` function. Do `?gausspr` and read the input arguments and the output. Also, go through my `KernLabDemo.R` carefully; you will need to understand it. Now, define your own square exponential kernel function (with parameters ℓ (`ell`) and σ_f (`sigmaf`)), evaluate it in the point $x = 1$, $x' = 2$, and use the `kernelMatrix` function to compute the covariance matrix $K(\mathbf{x}, \mathbf{x}_*)$ for the input vectors $\mathbf{x} = (1, 3, 4)^T$ and $\mathbf{x}_* = (2, 3, 4)^T$.
- (b) Consider first the model:

$$\begin{aligned} temp &= f(time) + \varepsilon \quad \varepsilon \sim N(0, \sigma_n^2) \\ f &\sim GP[0, k(time, time')] \end{aligned}$$

Let σ_n^2 be the residual variance from a simple quadratic regression fit (using the `lm()` function in R). Estimate the above Gaussian process regression model using the squared exponential function from 2a) with $\sigma_f = 20$ and $\ell = 0.2$. Use the `predict` function to compute the posterior mean at every data point in the training datasets. Make a scatterplot of the data and superimpose the posterior mean of f as a curve (use `type="l"` in the plot function). Play around with different values on σ_f and ℓ (no need to write this in the report though).

- (c) **kernlab** can compute the posterior variance of f , but I suspect a bug in the code (I get weird results). Do your own computations for the posterior variance of f (hint: Algorithm 2.1 in RW), and plot 95% (pointwise) posterior probability bands for f . Use $\sigma_f = 20$ and $\ell = 0.2$. Superimpose those bands on the figure with the posterior mean in 2b).
- (d) Consider now the model

$$\begin{aligned} \text{temp} &= f(\text{day}) + \varepsilon \quad \varepsilon \sim N(0, \sigma_n^2) \\ f &\sim GP[0, k(\text{day}, \text{day}')] \end{aligned}$$

Estimate the model using the squared exponential function from 2a) with $\sigma_f = 20$ and $\ell = 6 \cdot 0.2 = 1.2$. (I multiplied ℓ by 6 compared to when you used **time** as input variable since **kernlab** automatically standardizes the data which makes the distance between points larger for **day** compared to **time**). Superimpose the posterior mean from this model on the fit (posterior mean) from the model with **time** using $\sigma_f = 20, \ell = 0.2$. Note that this plot should also have the **time** variable on the horizontal axis. Compare the results from using **time** to the ones with **day**. What are the pros and cons of each model?

- (e) Now implement a generalization of the periodic kernel given in my slides from Lecture 2 of the GP topic (Slide 6)

$$k(x, x') = \sigma_f^2 \exp\left(-\frac{2 \sin^2(\pi |x - x'| / d)}{\ell_1^2}\right) \times \exp\left(-\frac{1}{2} \frac{|x - x'|^2}{\ell_2^2}\right).$$

Note that we have two different length scales here, and ℓ_2 controls the correlation between the same day in different years (ℓ_2). So this kernel has four hyperparameters σ_f, ℓ_1, ℓ_2 and the period d . Estimate the GP model using the **time** variable with this kernel with hyperparameters $\sigma_f = 20, \ell_1 = 1, \ell_2 = 10$ and $d = 365/\text{sd}(\text{time})$. The reason for the rather strange period here is that **kernlab** standardized inputs to have standard deviation of 1. Compare the fit to the previous two models (with $\sigma_f = 20, \ell = 0.2$). Discuss the results.

3. Gaussian process classification using the kernlab package. Download the banknote fraud data:

```
data <- read.csv('https://github.com/STIMALiU/AdvMLCourse/raw/master/GaussianProcess/Code/banknoteFraud.csv',
header=FALSE, sep=',')
names(data) <- c("varWave", "skewWave", "kurtWave", "entropyWave", "fraud")
data[,5] <- as.factor(data[,5])
```

You can read about this dataset here. Choose 1000 observations as training data using the following command (i.e. use the vector **SelectTraining** to subset the training observations)

```
set.seed(111); SelectTraining <- sample(1:dim(data)[1], size = 1000, replace = FALSE)
```

- (a) Use **kernlab** to fit a Gaussian process classification model for **fraud** on the training data, using **kernlab**. Use **kernlab**'s the default kernel and hyperparameters. Start with using only the first two covariates **varWave** and **skewWave** in the model. Plot contours of the prediction probabilities over a suitable grid of values for **varWave** and **skewWave**. Overlay the training data for **fraud** = 1 (as blue points) and **fraud** = 0 (as red points). You can take a lot of code for this from my **KernLabDemo.R**. Compute the confusion matrix for the classifier and its accuracy.
- (b) Using the estimated model from 3a), make predictions for the testset. Compute the accuracy.
- (c) Train a model using all four covariates. Make predictions on the test and compare the accuracy to the model with only two covariates.

Good luck!

- Mattias